Alberto Franzin, Fabio Palese

Sistemi Intelligenti

January 15th, 2013

Introduction

Introduction

Bayesian networks
Definition
Naive Bayes
Naive Bayes for spam classification

SpamBayes Implementation

Tests

Results and Conclusions

► Bayes rule:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

defines the *a posteriori* probability of event *A*, knowing the event *B* has already occurred.

► Bayes rule:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

defines the *a posteriori* probability of event *A*, knowing the event *B* has already occurred.

► In other words, we can estimate the probability of an hypothesis, given that we know the consequences.

► Bayes rule:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

defines the *a posteriori* probability of event *A*, knowing the event *B* has already occurred.

- ► In other words, we can estimate the probability of an hypothesis, given that we know the consequences.
- ► This has led to two different interpretations of the theorem.

► Bayes rule:

Introduction

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

► Bayes rule:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

ightharpoonup P(A|B) is the *a posteriori* probability

► Bayes rule:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

- ▶ P(A|B) is the *a posteriori* probability
- ▶ P(B|A) is the *likelihood*

► Bayes rule:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

- ▶ P(A|B) is the *a posteriori* probability
- ▶ P(B|A) is the *likelihood*
- ▶ P(B|A)P(A) is the *prior* probability

► Bayes rule:

Introduction

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

- ▶ P(A|B) is the *a posteriori* probability
- ▶ P(B|A) is the *likelihood*
- ▶ P(B|A)P(A) is the *prior* probability
- ► $P(B) = \sum_{a \in A} P(B|A = a)P(A = a)$ is the *total* probability

Frequentists vs. Bayesians



from http://xkcd.com/1132, see also http://en.wikipedia.org/wiki/Sunrise_problem

Frequentists vs. Bayesians



from http://xkcd.com/1132, see also http://en.wikipedia.org/wiki/Sunrise_problem

The frequentist relies on the theoretical probability of the events.

Frequentists vs. Bayesians



from http://xkcd.com/1132, see also http://en.wikipedia.org/wiki/Sunrise_problem

The bayesian observes the past events occurred, and adapts the probability accordingly.

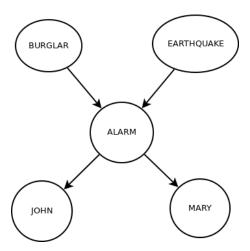
A Bayes network is a way to describe causal relationships between events.

- ► Nodes = events
- ► (Directed) Edges = causal relationship

A Bayes network is a way to describe causal relationships between events.

- ► Nodes = events
- ► (Directed) Edges = causal relationship
- ► Two nodes are connected by an edge: the child of an arc is influenced by its ancestor in a probabilistic way

AN EXAMPLE



CONDITIONAL INDEPENDENCE

▶ If

$$P(A|B,C) = P(A|B)$$

then we say that *B* and *C* are *conditionally independent*.

CONDITIONAL INDEPENDENCE

▶ If

$$P(A|B,C) = P(A|B)$$

then we say that *B* and *C* are conditionally independent.

▶ Note that conditional independence \neq independence

CONDITIONAL INDEPENDENCE

► If

$$P(A|B,C) = P(A|B)$$

then we say that *B* and *C* are *conditionally independent*.

- ▶ Note that conditional independence \neq independence
- Explaining away: if we know that one possible cause of the event has happened, this may explain away the event, being all the other causes less probable once we know the one that happened.

NAIVE BAYES

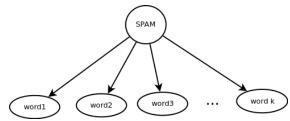
► Computing all the probabilities in a Bayesian network requires exponential time. We introduce the assumption of independence among variables.

NAIVE BAYES

- ► Computing all the probabilities in a Bayesian network requires exponential time. We introduce the assumption of independence among variables.
- ► It is called *naive*, since it's often unrealistic, but it yields good results.

NAIVE BAYES

- ► Computing all the probabilities in a Bayesian network requires exponential time. We introduce the assumption of independence among variables.
- ► It is called *naive*, since it's often unrealistic, but it yields good results.
- ► In spam classification:



NAIVE BAYES FOR SPAM CLASSIFICATION

We want to build a classifier that distinguishes good mails from undesired mails:

- ▶ good mails: *ham*
- ▶ undesired mails: *spam*

ALGORITHM

Algorithm:

► Training

ALGORITHM

Algorithm:

- ► Training
- ► Validation

ALGORITHM

Algorithm:

- ► Training
- ► Validation
- ► Testing

Formulas:

► For each word:

$$P_{word|spam} = \frac{\text{\# occurrences of word in spam mails}}{\text{\# total occurrences of word}}$$

Tests

NAIVE BAYES FOR SPAM CLASSIFICATION

Formulas:

► For each word:

$$P_{word|spam} = \frac{\text{# occurrences of word in spam mails}}{\text{# total occurrences of word}}$$

► Final for spam is:

$$P_{spam} = \prod_{words \in mail} P_{spam|word}$$

Formulas:

► For each word:

$$P_{word|spam} = \frac{\text{# occurrences of word in spam mails}}{\text{# total occurrences of word}}$$

SpamBayes

► Final for spam is:

$$P_{spam} = \prod_{words \in mail} P_{spam|word}$$

▶ same for ham

Formulas:

► For each word:

$$P_{word|spam} = \frac{\text{\# occurrences of word in spam mails}}{\text{\# total occurrences of word}}$$

SpamBayes

► Final for spam is:

$$P_{spam} = \prod_{words \in mail} P_{spam|word}$$

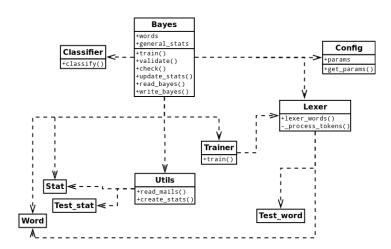
- same for ham
- ► Outcome is the class that maximizes the probability of belonging to that class.

SPAMBAYES

- ► SpamBayes: applying Naive Bayes to spam classification
- ► Python with Ply and BeautifulSoup
- ► dataset: SpamAssassin archive
- ► Code and documentation available at http: //code.google.com/p/sist-int-2012project/

SPAMBAYES

Introduction



NOTES ON IMPLEMENTATION

► Smoothing:

$$P_{word|spam} = \frac{\text{\# occurrences of word in spam mails} + k}{\text{\# total occurrences of the word} + |C| \times k}$$

NOTES ON IMPLEMENTATION

► Smoothing:

$$P_{word|spam} = \frac{\text{\# occurrences of word in spam mails} + k}{\text{\# total occurrences of the word} + |C| \times k}$$

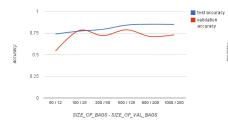
► Calculations can be simplified: some words bring little contribution to the mail status

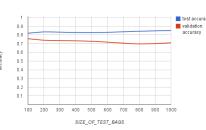
► Smoothing:

$$P_{word|spam} = \frac{\text{\# occurrences of word in spam mails} + k}{\text{\# total occurrences of the word} + |C| \times k}$$

- ► Calculations can be simplified: some words bring little contribution to the mail status
- ► Several parameters to be tuned: we describe the more relevant ones

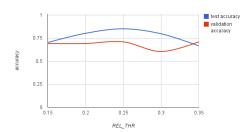
Size of training/validation/test sets:



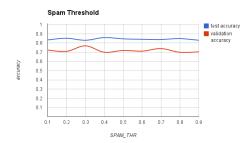


Tests

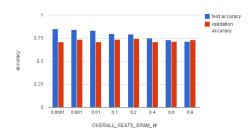
Relevance threshold:



"Spamicity" threshold:



Feature statistic threshold:



► About the dataset:

- ► About the dataset:
 - ► General features

- ► About the dataset:
 - ► General features
 - ► Single words

- ► About the dataset:
 - ► General features
 - ► Single words
- ► About the classification:

- ► About the dataset:
 - ▶ General features
 - ► Single words
- ► About the classification:
 - ► Accuracy

- About the dataset:
 - ► General features
 - ► Single words
- ► About the classification:
 - Accuracy
 - ► How to improve?